

Subham Joshi

Data Analyst · ML Evaluation · Americas-shift compatible (IST, 4 hrs PST overlap)

joshisubham442@gmail.com | linkedin.com/in/subham-joshi-b25778280 | github.com/Zolow-kuni/data-portfolio |
Dehradun, India · Remote · Open to PST-overlap hours

PROFESSIONAL SUMMARY

Data Analyst with ~1.5 years of hands-on experience in Python (Pandas, NumPy, Scikit-learn), SQL, and statistical analysis - with a focus on ML model evaluation, anomaly detection, and data quality validation. **Google Advanced Data Analytics Certified** (full 7-of-7 Professional Certificate, May 2026) covering EDA, statistics, regression, supervised learning (Random Forest, XGBoost, Decision Trees), and model evaluation (AUC/ROC, confusion matrices, feature importance). Built and evaluated 3 ML models on real production-scale datasets: Logistic Regression (AUC 0.837), Random Forest (0.824), and Decision Tree (0.781). Backed by 9 end-to-end portfolio projects across Python, SQL, R, and Power BI on 1.2M+ records. **Americas-shift compatible** - available for 4+ hours of PST overlap daily.

CORE COMPETENCIES

ML Model Evaluation (AUC/ROC/F1)	Statistical Analysis & Hypothesis Testing	Python (Pandas · NumPy · Scikit-learn)
SQL (Window Functions · CTEs)	Anomaly & Outlier Detection	Data Quality & Validation
EDA & Pattern Recognition	R (randomForest · caret · pROC)	Reproducible Analysis Workflows
Large Dataset Analysis (1.2M+ records)	Power BI · Tableau · Plotly	Documented, Clean Code (GitHub)

WORK EXPERIENCE

Logistics Integrators Pvt. Ltd.

Mar 2025 - Mar 2026 · Dehradun, India

Data Analyst (Management Trainee Program)

- Built **SQL-based statistical anomaly detection** and threshold-based alerting on operational and sales datasets - using Z-score and IQR logic to flag outliers before SLA impact; analogous to diagnosing ML model failure modes and edge cases in production
- Owned **data quality and integrity workflows**: validation, reconciliation, and ETL audit across reporting systems - eliminating discrepancies and maintaining analysis-ready datasets
- Conducted **root-cause analysis (RCA)** on performance anomalies - decomposed multi-variable incidents, traced causes through data lineage, and partnered with 3-5 cross-functional teams to validate findings and ship corrective actions
- Automated recurring analytical workflows in **Python (Pandas, NumPy)**, producing reproducible, well-documented reporting pipelines with audit logging and timestamped outputs
- Led **conversion and funnel analysis** on shipment journeys - computing drop-off rates, stage-by-stage throughput, and cohort-level performance patterns

Vrun Minerals Pvt. Ltd.

Feb 2024 - Jul 2024 · Dehradun, India

Operations Analyst (Intern)

- Cleaned and validated distributor and sales datasets using Python (Pandas) and Excel; built deviation-highlighting dashboards on inventory and sales trends
- Performed inventory reconciliation and data accuracy checks; documented workflows and metric definitions for non-technical stakeholders

PROJECTS - ML EVALUATION & ANALYTICS

Customer Churn Prediction - ML Model Evaluation

ML Evaluation · R

R · tidyverse · randomForest · caret · pROC · ggplot2

Trained and evaluated 3 ML models on 7,043 Telco customer records. **Logistic Regression: AUC 0.837** (best); Random Forest: 0.824; Decision Tree: 0.781. Produced 8 EDA plots, ROC curves per model, confusion matrices, and feature importance charts. Full model comparison workflow with reproducible R code. [\[GitHub\]](#)

Fraud & Anomaly Detection System Anomaly Detection · Python

Python · Pandas · NumPy · SciPy · Scikit-learn

Applied **Z-score**, **IQR**, and **rule-based risk scoring (0-100)** to 283,726 real European credit card transactions; flagged 2,838 suspicious records at a 0.17% fraud rate on a 284,807-row dataset. Includes threshold calibration, false-positive analysis, and explainable scoring output. [\[GitHub\]](#)

Google ADAC Capstone - Automating EDA (408K NYC Taxi Trips) EDA · PACE Framework

Python · Pandas · NumPy · Matplotlib · Seaborn · Tableau

Full EDA on 408,294 NYC TLC Yellow Taxi trips (18 columns); engineered 5 features, applied IQR outlier detection, flagged ~0.5% zero-distance anomalies. Built 4 Python visualisations + Tableau scatter plot; documented findings using PACE framework and executive summary - demonstrating structured analytical artifact creation. [\[GitHub\]](#)

Financial Data Integrity Checker Data Quality · Validation

Python · Pandas · NumPy · Matplotlib

10-rule validation engine on 700 Microsoft financial records; all rules passed, 14 profit outliers flagged via 3-sigma detection, completeness scored 91-100% per column. Reproducible, documented validation pipeline. [\[GitHub\]](#)

SQL Retail Analysis - Cohort, RFM & Statistical Reporting SQL · Statistics

SQL · SQLite · Python

Loaded 805,549 rows from 525,461 transactions; wrote 15 advanced queries covering window functions, CTEs, RFM segmentation, cohort retention, and MoM revenue growth via LAG() - across 5,878 customers in 41 countries. [\[GitHub\]](#)

EDUCATION

MBA / PGPM - Operations Management

2025 · CGPA 6.83

ICFAI Business School, Mumbai

Focus: data-driven decision making, process optimization, performance analytics.

Bachelor of Computer Applications (BCA)

2022 · 72.68%

Shree Guru Ram Rai I.T.S., Dehradun

Programming (Java, Python), relational databases (SQL), analytical problem solving.

CERTIFICATIONS

 Google Advanced Data Analytics Professional Certificate (7/7) - EDA, Statistics, Regression, ML, Capstone	Google / Coursera	May 2026
The Nuts and Bolts of Machine Learning - supervised learning, Random Forest, XGBoost, model tuning	Google / Coursera	May 2026
The Power of Statistics - hypothesis testing, probability distributions, inference	Google / Coursera	May 2026
SQL and PostgreSQL: The Complete Developer's Guide (22 hrs)	Udemy	May 2026
Python for Data Science & Data Analysis Using Python	IBM / Credly	2026

SKILLS

ML Evaluation: AUC/ROC Analysis, Confusion Matrix, Feature Importance, Model Comparison, Supervised Learning (Logistic Regression, Random Forest, Decision Tree, XGBoost), Outlier Detection (Z-score, IQR), Statistical Inference, Hypothesis Testing

Python: Pandas, NumPy, SciPy, Scikit-learn, Matplotlib, Seaborn, openpyxl, Streamlit, Plotly - clean, documented, reproducible pipelines

SQL: PostgreSQL, SQLite - joins, window functions, CTEs, aggregations, LAG/LEAD, RFM, cohort analysis

R: tidyverse, randomForest, caret, pROC, ggplot2 - full ML pipeline including EDA, training, evaluation, ROC curves

BI & Reporting: Power BI (DAX), Tableau, dbt (Fundamentals cert), ETL Validation, Data Quality Workflows